

Twitter Topic Recommendation with a Deep Learning Model

Anil Trikha, Pengshuai Shi.

I. INTRODUCTION

TWITTER is an online micro-blogging social media platform where users post messages of at most 140 characters in length. Users have developed a convention whereby the topic to which a tweet relates is encoded in the text of the message as a word beginning with a hash symbol. However, only 14.6% of tweets contain such a hashtag.

In this project we attempt to reproduce the modeling approach of Li et al [1], by training a deep neural network on a set of tweets containing hashtags in order to generate a hashtag for tweets that do not contain one. Classifying tweets in this way would potentially enable user interest modeling, recommendation services, and online advertising.

II. INDEPENDENT VARIABLES

The independent variables in the model are the words and sentences of which a tweet is composed. Approaches such as TF-IDF, kNN, and Naive-Bayes would learn a model using these variables directly. However, there are bottlenecks to such models because they do not make use of the semantic information inherent in words and their composition into sentences. We employ the word2vec technique introduced by Mikolov et al [2], 2013 to capture this meaning.

III. LITERATURE REVIEW

Much recent work in natural language processing (NLP) involves embedding words as vectors using neural networks (Mikolov et al, 2013 [2]) and performing classification on top of that (Collobert et al., 2011 [3]). We use the conclusions of such work to inform how we build an in-house classifier for the topic of the tweets we collect.

IV. PROPOSED APPROACH

We propose to follow the approach of Li et al [1], by first representing each word in the tweet as a numeric vector of multiple dimensions. The vectors are generated from an independent corpus of words using a tool, such as *gensim*, and then the tweet is converted to sets of these vectors, with one set per sentence. The technique, called word2vec, was originated by Mikolov et al [2] and has the property that similar words appear close to each other in the embedded space and thus capture some of the semantics of the word. A convolutional neural network (CNN) is trained with the word vectors of each sentence of a tweet as the input layer. This yields a set of sentence representations that are fed into an LSTM-RNN to yield a layer representing the tweet which is then fed into a softmax layer whose output is the probability of one of K different hashtags used as class labels.

V. DATA PREPARATION

A. Data Collection

The dataset was collected from Twitter using the *R searchTwitter* function. Five topics were selected: *basketball*, *hockey*, *baseball*, *volleyball*, *tennis* and a hashtag was added during collection, e.g. *#basketball*. The reason why these sports topics were chosen is because they are major sports so that the number of records collected would be sufficiently large for the purpose of the project. For instance, *golf* was one of the sports was originally selected, but there were too few data points and thus it was removed. *Football* also used to be one of the candidates, but due to the ambiguity of football and soccer in different regions it was also removed. Only five classes were collected since more classes would have slowed training. Table I below shows the final number of tweets collected,

Topic(#)	basketball	hockey	baseball	tennis	volleyball
Number of tweets	10,000	10,000	10,000	10,000	4,755

TABLE I: Collected Tweets

B. Data Pre-processing

Most of the conventional text pre-processing methods are used, with two major differences in this project. Firstly, all hashtags for each of the five sports are removed, since they are the labels for each of the tweets. But this is exactly the label that the model tries to predict so the model would learn nothing but would appear to be performing really well (high accuracy by just looking at the hashtags). Secondly, the user names are not removed as they could provide clues about which sport topic the tweet belongs to, e.g. “@KingJames” is the username of the basketball star player LeBron James and tweets involving this username are most likely about basketball. Thus the major pre-processing steps are as follows:

- 1) Replace *http* links with whitespace.
- 2) Replace hashtags {*#basketball*, *#hockey*, *#baseball*, *#volleyball*, *#tennis*} with whitespace.
- 3) Replace newline characters such as *\r*, *\n* and *\r\n* with whitespace.
- 4) Replace special symbols such as *#*, *@*, *%*, *&*, *\$*, digits and punctuations with whitespace.
- 5) Remove ‘RT’s (retweets).
- 6) Combine multiple whitespaces to one.
- 7) Remove duplicated tweets.

The order of 2), 3) and 4) matters since if 4) is executed first then 2) and 3) would not work correctly. 1), 5) and 7) are

motivated by the same reason as they do not contribute and are noise to topic classification.

VI. EXPERIMENTS WITH DESCRIPTIVE ANALYTICS

A. Benchmark Experiments

To compare the performance of our model to existing techniques, we ran several experiments to arrive at a benchmark based on these other methods.

In our first pass, we did not remove hashtags for the sports of interest from the tweets in the data pre-processing and cleaning step. We generated a document-term matrix, using the *RTextTools* R library, where each tweet was considered as a document. Frequently occurring terms were chosen as the features to be used for a classifier. To arrive at a reasonable number of features we varied a parameter for removing sparse terms when creating the matrix. With hashtags present, a Naive-Bayes classifier achieved 91% accuracy as per Table II.

predicted \ actual	baseball	basketball	hockey	tennis	volleyball
baseball	2765	5	1	5	142
basketball	4	2679	5	7	246
hockey	5	6	2718	0	353
tennis	8	17	0	2602	398
volleyball	9	20	1	4	1427

TABLE II: Predicted sports vs actual with hashtags kept and $removeSparseTerms = 0.9$

Having realized that the hashtags are essentially the class labels, we understood that their presence in the tweet text causes the classifier to simply recognize the hashtag when training. This results in inflated accuracy and does not reflect the fact that our attempt is to classify tweets that may not include the hashtag word directly. So, in our next pass, we cleaned the tweet text of hashtags, created the document-term matrix to yield more terms, and reapplied the Naive-Bayes classifier. The new accuracy was much lower at only 39% as per Table III.

predicted \ actual	baseball	basketball	hockey	tennis	volleyball
baseball	1005	79	42	19	1773
basketball	89	1494	179	52	1127
hockey	50	60	504	13	2455
tennis	34	84	77	830	2000
volleyball	16	17	12	4	1412

TABLE III: Predicted sports vs actual with hashtags removed and $removeSparseTerms = 0.985$

Since the new accuracy was much lower, we sought to verify the assumptions of the Naive-Bayes model. As the order of words in a tweet is important to understanding its content, the independence assumption may not hold. Also the features are so sparse in the columns of the document-term matrix that the assumption of a Gaussian distribution does not hold. Figure 1 shows the distribution of the term *nba* from the previous experiment.

To explore whether we could achieve higher accuracy with other classification techniques, we reduced the data set to 500

Distribution of NBA term in Document-Term Matrix




Fig. 1: Non-Gaussian distribution of NBA in document-term matrix

tweets per class for a total of 2500, using 70% for training and 30% for test. This reduction enabled our experiments to complete in a reasonable amount of time. For Naive-Bayes we noted a revised accuracy of 38%. Random Forest achieved an accuracy of 27% as per Table IV and SVM achieved an accuracy of 24% as per Table V.

predicted \ actual	baseball	basketball	hockey	tennis	volleyball
baseball	12	51	16	43	40
basketball	2	52	12	27	33
hockey	14	52	11	47	46
tennis	6	35	5	65	18
volleyball	7	54	22	20	60

TABLE IV: Actual vs predicted sports as classified by a Random Forest model trained on a reduced dataset

predicted \ actual	baseball	basketball	hockey	tennis	volleyball
baseball	12	66	23	33	28
basketball	5	56	13	24	28
hockey	18	65	12	37	38
tennis	15	43	4	49	18
volleyball	11	60	21	18	53

TABLE V: Actual vs predicted sports as classified by SVM with a radial kernel and trained on a reduced dataset

The SVM classification accuracy for various kernels is plotted in Figure 2.

Clearly, our experiments show that after cleaning tweets of the hashtag labels for the sports under consideration, the best test accuracy attainable with conventional methods is 39%. We now describe our LSTM model and demonstrate its superiority on this task.

VII. PREDICTIVE MODELING

A. Feature Engineering/extraction

Each data point in the raw input data set is a tweet that consists of a list of words, which we convert to vector representation and use as the input features to our model. As we have seen previously, using term frequency as the features of each tweet has its limitations since it does not capture




Fig. 2: SVM Test Accuracy as achieved by various kernels on the reduced dataset

semantic information very well. Word2vec techniques, more specifically the *Skip-Gram* model are used here to convert a string of words to a vector representation of features. See Figure 3.




Fig. 3: The skip-gram model architecture [1].

The skip-gram model utilizes a shallow one-layer neural network to learn the representation of words such that the likelihood of obtaining surrounding words from a given word is maximized. In this project, we mapped each unique word into a 600-d low dimensional (compared to one-hot encoding) vector space, i.e. the number of features used in our project is 600. We could convert the data to even higher dimensional feature spaces, but that would slow the training and is not necessary as the model is already performing very well.

B. Model Design

In Li et al [1], they first start with a one-hot word representation, pass it through word2vec layers and form a distributed word representation. Then a convolutional neural network (CNN) is applied to each sentence in a tweet to extract local semantic information, and finally the output of the CNN is fed to Long Short-Term Memory (LSTM) layer.

Due to time constraints, in this project our model architecture is different from theirs. We directly use an already trained word2vec model on Wikipedia Text8 corpus to encode each word in the input tweets, instead of training the word2vec representation ourselves. There is no sentence composition, that is we do not apply CNN on each sentence in a tweet. Instead, the memory cell in the LSTM layer is directly connected to each vector represented word in our model, and the final output is the predicted topic of that tweet, see Figure 4. We do not use a softmax layer for the classification.

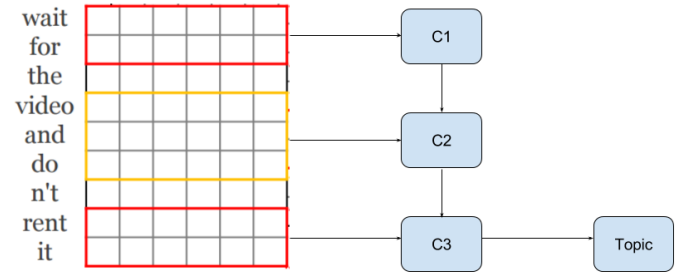


Fig. 4: Each word in sentence connects to a LSTM cell

Each of the memory cells contains 4 major operations:

- Forget gate: Provide a control over how much information comes through from previous cell.
- Input node: Activation from input data at current time step.
- Internal gate: Update and provide different candidate.
- Output gate: Decide what to output and also what gets passed to next memory cell.

C. Build Model

1) *Train*: Combine all sports tweets collected into one data set, with total 44,755 number of tweets. Train the model with a stratified 5-fold cross validation, each with 75% training set and 25% validation set.

The number of hidden neurons is 300, the number of memory cells is 23 which is our sentence length (number of words), the dimension of each word is 600 as mentioned previously (so the total input data dimension is 23x600). We use a batch size of 50 and 20 epochs to train our model.

2) *Validation*: Predictions are made on the 25% validation set from previous split, 5 times in total.

3) *Result*: The evaluation metric of the result on the validation set is *harsh accuracy* which is the mean of the accuracy per class.

Table VI shows the result of the first cross validation.

predicted \ actual	baseball	basketball	hockey	tennis	volleyball	Accuracy
baseball	2335	35	28	64	27	0.9381
basketball	53	2356	52	78	42	0.9128
hockey	24	48	2289	56	16	0.9408
tennis	44	56	35	2342	32	0.9334
volleyball	23	42	28	72	1012	0.8598

TABLE VI: LSTM prediction result of 1st cross validation: confusion matrix and accuracy per class

Finally, we show the complete result of accuracy per class and harsh accuracy for each fold in the cross validation in Figure 5:




Fig. 5: Accuracy per class and overall harsh accuracy for each validation set

We see that it achieves a remarkable result on this 5-class sports classification task, with high harsh accuracy and low variance. Although the model's prediction on *volleyball* does not perform as well as the other topics, this is probably due to the relatively small size of the volleyball data set compared to the other collected sports.

VIII. CONCLUSION

In this project, we developed a deep learning system to simulate the assignment of a sports topic hashtag to tweets that do not have one. We limited the scope of our work to classifying a tweet as one of five popular sports. The deep learning architecture chosen was an LSTM neural network fed with input consisting of the words of a tweet represented as vectors. The output of the LSTM was used directly for the class probability. The system achieved a test accuracy of 92%.

To compare our results to more conventional techniques, we performed a series of benchmark experiments on a data set that was reduced in size. These experiments used the TF-IDF method to obtain a document-term matrix before applying a classifier. Naive-Bayes, Random Forest, and SVM achieved test accuracy scores of 38%, 27%, and 24% respectively. Clearly, LSTM, with word2vec representations as input, greatly outperforms them all.

For future work, the model could possibly be brought more in line with that of Li et al [1] by using a CNN

for the sentences of a tweet which feeds an LSTM with its representations. With more computing resources, the number of different sports used as class labels could also be increased.

REFERENCES

- [1] J. Li, H. Xu, X. He, J. Deng and X. Sun. Tweet Modeling with LSTM Recurrent Neural Networks for Hashtag Recommendation *International Joint Conference on Neural Networks (IJCNN)*, 2016
- [2] T. Mikolov, I. Sutskever, K. Chen, G. Corrado, J. Dean. 2013. Distributed Representations of Words and Phrases and their Compositionality. *In Proceedings of NIPS*, 2013.
- [3] R. Collobert, J. Weston, L. Bottou, M. Karlen, K. Kavukcuglu, P. Kuksa. 2011. Natural Language Processing (Almost) from Scratch. *Journal of Machine Learning Research* 12:24932537